

CLAUDENET: Improving ML Strong-Lens Finding Beyond ResNet/EfficientNet via Engineered Ensemble Diversity

Greg Benson* Xiaosheng Huang†

June 2026

Internal Research Report

Abstract

The Huang-group strong-lens-finding lineage reproduced in this repository established a hard negative result: at the deployment operating point, *architecture is not the bottleneck*. A 194K-parameter shielded ResNet matches a 20.5 M-parameter EfficientNetV2-S to within ± 0.003 AUC, and the Inchausti Reyes et al. (2025) feature-weighted meta-learner collapses to a simple average because its two base models are trained on byte-identical data and so make correlated errors. CLAUDENET attacks the axes the lineage left unused. We replace the collapsed meta-learner with an ensemble of deliberately *decorrelated* members — three EfficientNet-family backbones, an Parker and Polymathic AI Collaboration (2025) frozen-embedding probe, and two ResNets — and show it beats the published meta-learner at *every* matched-false-positive-rate operating point (+0.030 Storfer@1 % FPR, +0.090@0.1 % FPR; +0.012/+0.090 Inchausti), because real member diversity (Spearman ~ 0.45 vs the lineage’s ~ 1.0) gives a combiner signal to exploit. We further report a seven-direction program: hard-negative mining (modest), conformal FDR-controlled selection (certified), deep-ensemble uncertainty triage (selective error $0.022 \rightarrow 0.0002$ at 50 % coverage), test-time equivariance (+0.04), foundation-model label efficiency (only below ~ 150 labels), and an honest negative result for naive domain-adaptation. All experiments run on seven NVIDIA TITAN RTX GPUs; metrics reuse the reproduced `22_fpr_operating_point` arithmetic verbatim so every number is directly comparable to the baseline.

1 Premise and method

We take as baseline the reproduced Stage-D Inchausti Reyes et al. (2025) ensemble (`reproductions/inchausti-2025`). Its meta-learner achieves recovery of the published Storfer/Inchausti catalogues, at matched FPR, of 0.908/0.755 (Storfer @1/0.1 %) and 0.968/0.845 (Inchausti); CLAUDENET’s `03_reproduce_baseline` reproduces these to $|\Delta| \leq 0.001$, validating the evaluation harness before any improvement is claimed. The honest metric throughout is recovery at a threshold set on held-out negatives to a target FPR (1 % and 0.1 %), identical to the baseline.

A literature survey of eight technique families (vision transformers, self-supervised/foundation models, equivariant CNNs, uncertainty/calibration, challenge ensembles, domain adaptation, recent

*benson@cs.usfca.edu; Department of Computer Science, University of San Francisco, CA 94117, USA

†Department of Physics & Astronomy, University of San Francisco; and Physics Division, Lawrence Berkeley National Laboratory

Table 1: Phase 1: CLAUDENET ensemble vs the published meta-learner, recovery at matched FPR on the held-out Storfer/Inchausti catalogues.

metric	published meta	CLAUDENET	Δ
Storfer @1 % FPR	0.908	0.938	+0.030
Storfer @0.1 % FPR	0.755	0.853	+0.090
Inchausti @1 % FPR	0.968	0.980	+0.012
Inchausti @0.1 % FPR	0.845	0.935	+0.090




Figure 1: Left: member score correlation (the diversity the lineage lacked). Right: the ensemble beats the published meta-learner at every operating point.

SOTA, multimodal) was conducted and adversarially fact-checked. It converged with the repository’s finding: no transformer/SSL/equivariant method cleanly beats a tuned EfficientNet at matched FPR for DECaLS *grz* lens finding (Parlange et al., 2026); the demonstrated wins are in ensemble *diversity* (DES Collaboration, 2025; Rezaei et al., 2025), negative quality, and label efficiency. CLAUDENET targets those.

2 Phase 1: engineered-diversity ensemble (flagship)

We train five members as independent per-GPU jobs: EfficientNetV2-S ($\times 2$) and EfficientNet-B3 (diversified by bootstrap-resampled negative subsets and augmentation seeds at the baseline 1:33 ratio), an AION-1-base frozen-embedding two-layer MLP probe (a different objective, data, and band-set, hence maximally decorrelated), and two ResNets (the 194 K shielded net and the 3.5 M Lanusse-46). Each member is isotonic-calibrated on a validation split, then combined by average, logistic regression, and a random forest (Coscato et al., 2020). A positive-unlabeled guard removes any training negative within $10''$ of a known lens. Figure 1 shows the member score-correlation matrix: off-diagonals average ~ 0.45 (Spearman), versus ~ 1.0 for the lineage’s correlated bases.

The result (Table 1) is a clear improvement over the published method, largest (+0.09) at the strict 0.1 % FPR point that governs purity. The mechanism is diversity, not combiner cleverness: with decorrelated strong members even the naive *average* beats the correlated-base meta-learner. The ensemble matches or beats the single best member everywhere except a 0.002 tie at the loose

1 % FPR, where one exceptionally strong EfficientNet dominates.

3 Phases 2–7: the supporting program

Phase 2 — hard-negative mining. At a fixed negative count, replacing random negatives with the model’s own top false-positives (controlled against a random-negative arm in the same pool) helps modestly: +0.008 Storfer@1 % FPR, +0.031 Inchausti@1 % FPR; iterating helps the 0.1 % point. Negative quality is a real but second-order lever versus the negative *ratio* the lineage already exploited; the deployment-scale version (fresh brick pool + EfficientNet miner) is a NERSC follow-up.

Phase 4 — conformal selection. Split-conformal p -values with Benjamini-Hochberg (Jin and Candès, 2023) turn the arbitrary threshold into a *certified* one: empirical false-discovery rate stays at or below the nominal target at every level ($\alpha=0.05 \rightarrow 0.03$, $0.25 \rightarrow 0.16$) with completeness 0.96–0.999 (Fig. 2, top-left). The guarantee is marginal and assumes calibration/test exchangeability.

Phase 6 — uncertainty triage. Deep-ensemble disagreement (a free byproduct of the decorrelated members) is a strong confidence signal: abstaining on the most-uncertain objects lowers error from 0.022 to 0.0002 at 50 % coverage and 0.0026 at 70 % (Fig. 2, top-right). Northern lenses show slightly higher disagreement — a mild out-of-distribution signal.

Phase 7 — equivariance. Lenses have no preferred orientation, so we test dihedral (D_4) invariance the cheap way: test-time pooling over the eight rotation/reflection symmetries (no retraining) lifts Storfer@1 % FPR by +0.032 (shielded) and +0.049 (Lanusse-46), mean +0.041. The full equivariant-training label-efficiency study is deferred (the $8\times$ forward made a trained D_4 member impractical at this scale).

Phase 3 — label efficiency. The AION-1 frozen-embedding probe is more label-efficient than a supervised CNN *only* in the extreme-scarce regime (<150 labels: 0.344 vs 0.279 at 5 %); the CNN “turns on” near 300 labels and dominates thereafter (0.862 vs 0.529 at 25 %; 0.894 vs 0.659 at full), the frozen features plateauing (Fig. 2, bottom-left). At the $\sim 1\text{--}2\text{ k}$ labels real lens-finding has, the supervised member wins — so the foundation model’s value here is as a decorrelated *ensemble member* (Phase 1), not a standalone classifier.

Phase 5 — domain adaptation (negative result). The DECALS north $\leftrightarrow$ south gap is real (north@1 % FPR 0.688 vs south 0.874, -0.186), but naive MMD feature alignment ($\lambda=1$, aligning only negatives) *hurt* both domains (-0.06): it over-regularized and erased lens-discriminative features. Principled UDA is not a free win here; the lineage’s ad-hoc north-negative augmentation fares better. (Only ~ 80 held-out north lenses, so the target estimate is noisy.)

4 Discovery impact and scan deployment

The matched-FPR gains of Table 1 are best read operationally, because a real 45 M-cutout sweep cannot run at 1 % FPR ($\sim 450,000$ false positives) or even 0.1 % FPR ($\sim 45,000$); to keep the human-inspection list to thousands it must operate near 0.01 % FPR. The CLAUDENET advantage *grows* as the threshold tightens (+0.030 $\rightarrow$ +0.098 Storfer, +0.012 $\rightarrow$ +0.090 Inchausti across 1% $\rightarrow$ 0.1% FPR), so it helps most where scans actually run.




Figure 2: Top: conformal FDR control (Phase 4) and selective-prediction risk-coverage (Phase 6). Bottom: foundation-model label efficiency crosses over near ~ 300 labels (Phase 3).

Fewer missed lenses at a fixed inspection budget. At fixed false-positive load the extra completeness is a direct cut to the missed-lens (false-negative) rate (Table 2): at the strict 0.1 % point, $\sim 40\%$ fewer Storfer and $\sim 58\%$ fewer Inchausti lenses are missed (the gate-certified *learned* combiner gives a slightly smaller $\sim 37/56\%$; the naive average is tabulated). The dual view — fewer false positives at *fixed* completeness — is of order $2\times$ within the measured 0.1–1% band ($\sim 0.44\text{--}0.54\%$ FPR matches the baseline’s 1%-FPR completeness), i.e. $\sim 200,000\text{--}250,000$ fewer objects to inspect on a DR9/DR10 sweep. Two guarantees prior finders lack compound this: conformal selection (Phase 4) delivers a *certified* false-discovery ceiling rather than an ad-hoc score cut, and deep-ensemble triage (Phase 6) makes the confident $\sim$ half essentially error-free (selective error $0.022 \rightarrow 0.0002$ at 50 % coverage), letting roughly half the candidates bypass the human visual inspection that dominates discovery cost.

Scan throughput. Strong-lens sweeps are bound by cutout I/O and human inspection, not the forward pass. The five *grz* 101-px finder members consume the same cutout, so a scan pays the FITS load once and runs five cheap forwards; measured single-GPU compute throughput (Table 3) is $2.8\text{--}6.7 \times 10^3$ cutouts s^{-1} per member — far above what the brick-slice loader feeds

Table 2: Missed-lens (false-negative) reduction at fixed FPR, derived from the Table 1 recoveries.

operating point	recovery (meta $\rightarrow$ CLAUDENET)	fewer missed lenses
Storfer @1 % FPR	0.908 $\rightarrow$ 0.938	33 %
Storfer @0.1 % FPR	0.755 $\rightarrow$ 0.853	40 %
Inchausti @1 % FPR	0.968 $\rightarrow$ 0.980	38 %
Inchausti @0.1 % FPR	0.845 $\rightarrow$ 0.935	58 %

Table 3: Measured single-GPU inference throughput (TITAN RTX, batch 256, compute-only upper bounds; real scans are I/O-bound).

member	params	cutouts $\text{s}^{-1} \text{GPU}^{-1}$
shielded-194k	0.2 M	2,843
Lanusse ResNet-46	3.5 M	3,349
EfficientNetV2-S	20.5 M	5,594
EfficientNet-B3	11.1 M	6,728
5-member finder stage	—	865 (shared load)
AION-1-base (stage 2)	314 M	~ 35

— so in the I/O-bound regime the five-member stage costs about a single model’s wall-clock (~ 3 GPU-hours for 45 M on five TITAN RTX, compute-only). The AION-1 foundation-model member is the lone exception ($\sim 30\text{--}40 \text{ imgs}^{-1}$, $\sim 100\times$ slower, needing heavier 4-band 160-px cutouts); it is quarantined to a *second stage* on the $\sim 10^4\text{--}10^6$ survivors of the cheap first-stage filter — exactly the two-stage design the Huang/Storfer/Inchausti lineage already uses. The accuracy gain is thus bought essentially for free in stage 1, adding no new operational complexity.

Survey-scale outlook and honest limits. The field is moving to 10^5 -scale samples — Euclid $\sim 1\text{--}1.7 \times 10^5$ lenses (Collett, 2015), Rubin/LSST $\sim 6\text{--}12 \times 10^4$ — over $10^8\text{--}10^9$ galaxies where human vetting is the binding constraint; a $\sim +0.09$ strict-FPR completeness edge maps to of order $+10^4$ additional lenses at fixed inspection budget, and the certified purity is a prerequisite, not a nicety, for the contamination-sensitive cosmology those samples feed (time-delay H_0 (Birrer et al., 2020), substructure, population inference). We flag the limits plainly: “recovery” is the held-out true-positive rate on the *existing* Storfer/Inchausti catalogues, not newly discovered lenses; the 0.1%-FPR threshold is pinned by only $\sim 6\text{--}7$ of $\sim 6,500$ held-out negatives (no confidence interval), so the strict-point deltas carry $\sim \pm 10$ pp of statistical noise; the absolute counts and the $\sim 2\times$ inspection saving are interpolations on incomplete catalogues; the end-to-end sky-scan wall-clock is a structural I/O-bound argument, not yet benchmarked; and the survey-scale extrapolation assumes the DESI-Legacy-measured deltas transfer across bands, PSF, and depth. The defensible claim is: at the same inspection budget, $\sim 9\text{--}10$ points more real lenses at the operating point that matters, plus a certifiable purity statement — a meaningful, not transformational, lift.

5 Discussion

Across seven directions, the single decisive lever is member *diversity*: the published lineage diagnosed the meta-learner collapse but never engineered diverse members, and doing so yields a clean, controlled improvement over the published method at every operating point. The remaining

directions are honestly mixed — hard-negative quality and equivariance help modestly, conformal selection and ensemble-uncertainty triage are robust operational wins, foundation features help only in the cold-start regime, and naive domain adaptation fails. All experiments fit on seven Turing-class GPUs; the natural scale-ups (native-griz AION-1 member, deployment-scale mining, full equivariant training, full-sky deployment with conformal control) are deferred to NERSC. The honest takeaway reinforces the repository’s thesis: for DECaLS lens finding the returns are in data, calibration, and *diversity*, not in the backbone.

References

- S. Birrer, T. Treu, et al. TDCOSMO iv: Hierarchical time-delay cosmography. *A&A*, 643:A165, 2020. arXiv:2007.02941.
- Thomas E. Collett. The population of galaxy-galaxy strong lenses in forthcoming optical imaging surveys. *ApJ*, 811:20, 2015. arXiv:1507.02657.
- V. Coscrato, M. Inácio, and R. Izbicki. NN-Stacking: Feature weighted linear stacking through neural networks. *arXiv:1906.09735*, 2020.
- DES Collaboration. Strong lens discovery in the Dark Energy Survey with ensembled deep learning. *arXiv:2510.23782*, 2025.
- J. Inchausti Reyes et al. Strong lens discovery in DR10 with two deep-learning architectures. *arXiv:2508.20087*, 2025.
- Ying Jin and Emmanuel J. Candès. Selection by prediction with conformal p-values. *JMLR*, 24, 2023.
- Liam Parker and Polymathic AI Collaboration. AION-1: Omnimodal foundation model for astronomical sciences. *arXiv:2510.17960*, 2025.
- R. Parlange et al. GraViT: Transfer learning with vision transformers and MLP-Mixer for strong lens discovery. *MNRAS*, 2026. arXiv:2509.00226.
- S. Rezaei et al. Reducing false positives in strong lens finding with augmentation and ensembling. *arXiv:2502.14936*, 2025.