

Reproducing the Huang 2020 DECaLS Lens Finder from Paper + Public Code

Greg Benson Xiaosheng Huang

May 2026

commit: ef4c3c1

Phase 3a Internal Reproduction Report

Abstract

We re-implement the CMU DEEPLENS RESNET-46 architecture of [Lanusse et al. \(2018\)](#) in PyTorch and train it under the [Huang et al. \(2020\)](#) recipe on 929 NeuraLens L18-model lens candidates plus 4,908 random galaxies sampled from the DESI DR1 spectroscopic catalog as non-lenses. With no access to the original training code, training set, or trained weights, we reach a held-out test AUC of 0.9991 (validation AUC 0.9983) in 24.8 minutes on a single NVIDIA L4. This narrowly exceeds the published [Huang et al. \(2020\)](#) validation AUC of 0.98 — but the comparison is biased upward by two methodological substitutions (NeuraLens catalog versus the historical Master Lens Database + 6 publications; uniformly random non-lens galaxies versus the paper’s curated by-eye hard negatives). We discuss the substitutions and their effect in §7, and signpost Phase 3b — full DECaLS deployment to recover the published 335 new candidates — as the natural next step.

Contents

1	Introduction	2
2	Architecture	2
3	Training set	3
4	Hyperparameters	3
5	Training run	4
6	Results	5
7	Reproducibility caveats	5
8	Software and data availability	6
9	Conclusions	7
A	Per-epoch training history (compact)	7

1 Introduction

This report documents Phase 3a of an 8-phase reproduction roadmap covering Prof. Xiaosheng Huang’s strong-gravitational-lensing program at the University of San Francisco. Phase 1 (Huang et al. (2025); Foundry I HST observations + GIGA-Lens modeling) and Phase 2 (Hsu et al., 2025) have already been reproduced in companion tech-reports.¹ Phase 3 is the ResNet/EfficientNet lens-finder lineage: Huang et al. (2020, 2021) and successors (Storfer 2024, Inchausti 2025).

Phase 3a — this report — narrowly scopes Huang et al. (2020): the original 9,000 deg² DECaLS search that produced 335 new lens candidates using the CMU DEEPLens RESNET of Lanusse et al. (2018). The intent of Phase 3a is methodological validation: *can we re-implement the same architecture under the same hyperparameter recipe from the paper alone and reproduce the published AUC of 0.98?* Full deployment on a few million DEV/COMP-typed DECaLS galaxies to re-discover the 335 candidates is Phase 3b, deferred to a separate workstream.

The reproduction was performed without access to Huang’s own training code, training set, or trained weights. Once obtained through collaboration, Huang’s code will be the proper ground-truth comparison target.

2 Architecture

We adopt the CMU DEEPLens RESNET-46 of Lanusse et al. (2018) verbatim, which is the same architecture Huang et al. (2020) re-implemented in TensorFlow. The original is Theano/Lasagne; ours is PyTorch (Paszke et al. (2019); lives at `reproductions/huang-2020/01_lanusse_resnet.py`).

The architecture (figure 1) is one 7×7 convolutional stem followed by five stages of three pre-activated bottleneck residual blocks each (He et al., 2016). Each block has the form $1\times 1 \rightarrow 3\times 3 \rightarrow 1\times 1$ with ELU activation and batch normalisation. Channel progression $32 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 256 \rightarrow 512$; spatial downsampling by 2 at the first block of every stage except stage 1, achieved with strided convolutions in the bottleneck and in the projection shortcut. The head is an adaptive average pool, a single fully-connected layer to one logit, and a sigmoid. Total: 46 convolutional layers, ~ 3.5 M parameters at 3-channel input.




Figure 1: Lanusse 2018 / CMU DeepLens RESNET-46. The 5-stage pre-activated bottleneck stack with channel progression $32 \rightarrow 512$. Output spatial resolutions are shown for a 101×101 input (Huang et al. (2020) cutout size).

For the same architecture run at Lanusse et al. (2018)’s native 45×45 input, the spatial chain reduces to $45 \rightarrow 23 \rightarrow 12 \rightarrow 6 \rightarrow 3$ which matches their Figure 4. A smoke test in `01_lanusse_resnet.py` verifies a forward+backward pass at both input sizes produces a real-valued logit, gradient-finite, with parameter count 3,508,833.

¹See `reproductions/foundry-i/papers/main.tex` and `reproductions/hsu-2025/papers/main.tex`.

3 Training set

Positives. 949 lens candidates from the NeuraLens combined catalog², filtered to `ResNet Model == 'L18'` (Huang+2020 era), parsed from the DESI-RA±DEC naming convention. Of these, 929 have successful `ls-dr9` cutouts on disk.

Negatives. 5,000 random galaxies sampled from the DESI DR1 redshift catalog (`zall-pix-iron.fits`; Huang et al. (2025)) with the filter cascade `ZCAT_PRIMARY && ZWARN==0 && SPECTYPE!='STAR' && Z>0.05 && Dec ∈ [−68, +32.5]`, with a 10" exclusion zone around any positive. The expected lens-rate contamination is $\sim 10^{-4}$ per Huang et al. (2020) §3.1, so negligible against 5,000 rows. After cutout download, 4,908 negatives are available.

Cutouts. 101×101 pixel grz, pixscale 0.262"/pix (DECam native), FoV $\approx 26.5''$, pulled from legacysurvey.org/viewer/fits-cutout on the `ls-dr9` layer. Same field-of-view and band selection as Huang+2020 §3.2.

Split. Stratified 70:20:10 random shuffle by label, seed 2026. The resulting counts are shown in Table 1.

Table 1: Stratified 70:20:10 split counts.

split	label		total
	0 (non-lens)	1 (lens)	
train	3,436	650	4,086
validation	982	186	1,168
test	490	93	583
total	4,908	929	5,837

Class imbalance. Our negative:positive ratio is 5.3:1; Huang et al. (2020) used 21:1 (13,000 vs 613). This makes our task moderately easier; we discuss the implication in §7.

4 Hyperparameters

Unchanged from Lanusse et al. (2018) §3.4 and Huang et al. (2020) §3.2:

- Optimizer: ADAM (Kingma and Ba, 2015), default $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$.
- Initial learning rate 1×10^{-3} , divided by 10 every 40 epochs.
- Batch size 128.
- 120 epochs.
- Loss: binary cross-entropy with logits (numerically stable form).

²https://drive.google.com/file/d/1_KbEHWhl8LeeTyXpXkWFbLRxt6o42wBg

- Preprocessing: per-band mean subtraction + per-band standard deviation normalisation, computed once over 500 randomly-sampled training cutouts; outputs clipped to $\pm 250\sigma$.
- Augmentation (training-only): random rotation in $[-90, +90]^\circ$, random horizontal and vertical flips, random zoom in $[0.9, 1.0]$ followed by centre-crop back to 101×101 .

The only methodological deviation from the original is the rotation implementation: PyTorch’s `torchvision.transforms.functional.rotate` with `InterpolationMode.BILINEAR` versus [Lanusse et al. \(2018\)](#)’s spline interpolation. We did not observe a difference attributable to this in informal A/B tests.

5 Training run

Training was performed on a single NVIDIA L4 (24 GB) at `/raid/benson/git/agent-ic-lensing`, using the PyTorch 2.12 + `torchvision` 0.27 stack against CUDA 12.6 (venv at `/raid/benson/.venvs/lensfinder`). Wall-clock for the full 120 epochs including held-out test pass was 24.8 min. Figure 2 shows the loss and validation AUC trajectories.




Figure 2: Train and validation losses (left axis) and validation AUC (right axis) per epoch. Vertical dashed lines mark the two learning-rate decay points at epochs 40 and 80. Best validation AUC (highlighted) was 0.9983 at epoch 93.

Early-epoch validation AUC is noisy in the 0.18–0.97 band because each epoch contains only 32 batches (4,086 training samples / 128 batch size), so single-batch noise dominates the per-epoch eval. After the first learning-rate decay at epoch 40 the validation curve stabilises above 0.99 and the best checkpoint lands at epoch 93. The final 40 epochs at 10^{-5} provide a small additional refinement.

6 Results

The best validation-AUC checkpoint (epoch 93) was applied to the held-out test split (583 cutouts, 93 positives) producing the receiver operating characteristic in figure 3. The test AUC is 0.9991.




Figure 3: ROC on the held-out test set. The classifier achieves perfect ranking on the top half of the threshold sweep; the few mis-rankings near the upper-left corner are visible. See §7 for why this AUC is biased above the level we should expect from a fair out-of-distribution test.

The headline comparison against Huang et al. (2020) is in Table 2.

7 Reproducibility caveats

The 0.9991 test AUC narrowly exceeds the 0.98 published Huang et al. (2020) validation AUC, but a fair interpretation must account for two methodological substitutions.

(i) **The NeuraLens catalog mixes Huang’s training inputs with his discoveries.** Huang et al. (2020) trained on ~ 613 known lenses sourced from the Master Lens Database (Moustakas et al., 2012) plus six specific recent publications (Carrasco 2017; Diehl 2017; Pourrahmani 2018; Sonnenfeld 2018; Wong 2018; Jacobs 2017 & 2019). The 949 L18-model rows we use as positives are the *output* of the trained classifier as re-graded by the team: that is, approximately the 613 training inputs *plus* the 335 newly-discovered candidates, minus a handful of re-grading rejections. The NeuraLens release table does not tag rows by provenance, so we cannot split training-input from discovery-output positives.

Table 2: Headline comparison. The AUC gap is dominated by the two methodological substitutions in §7, not by the architecture or hyperparameter recipe.

	Huang+2020 published	Ours (Phase 3a)
AUC	0.98 (validation)	0.9991 (test), 0.9983 (val)
Positives	613 (MLD + 6 publications)	929 (NeuraLens L18)
Negatives	13,000 (by-eye hard)	4,908 (random DR1)
Imbalance	21:1	5.3:1
Hardware	3× Haswell @ NERSC Cori	1× NVIDIA L4
Wall-clock	17 h	24.8 min
Cutouts	101 × 101 grz, DR3/5	101 × 101 grz, DR9

Consequence: our random 70:20:10 split places a fraction of the 93 test positives at systems Huang himself trained on. Those systems tend to be visually obvious well-curated lenses that survived MLD-style review — substantially easier than typical sky candidates. The test AUC is therefore biased upward relative to a true out-of-sample evaluation.

(ii) **Uniformly random DR1 galaxies are easy negatives.** Huang et al. (2020) §3.1 mentions an explicit curated set of “by-eye” hard negatives: spiral galaxies, galaxy groups, galaxies of different colors, cosmic rays, unusual configurations, and data- reduction artefacts (their Figure 2). These are the cases where a naive classifier would produce false positives. We instead sampled negatives uniformly from the DESI DR1 fiber catalog, which is dominated by smooth elliptical and disk galaxies with no close companions — the easy end of the discrimination problem.

Consequence: our model is being asked an easier yes/no question than the paper’s bench. A faithful reproduction at the same difficulty level needs the by-eye hard-negative set; that’s Phase 3b work.

Honest take. Phase 3a establishes that the *architecture and training recipe* can be re-implemented to a working classifier from the published methods alone. It does *not* establish parity with the paper’s bench-marked performance. The 0.9991 test AUC should be read as a lower bound on what is achievable, not as a like-for-like improvement over 0.98.

A proper closure of the Phase 3a → paper gap requires either (a) obtaining Huang’s training set and code via collaboration, or (b) assembling an authentic 613-row MLD+publications training-positive set *and* a 13,000-row by-eye-hard-negative set, then re-running.

8 Software and data availability

- Architecture and training: `reproductions/huang-2020/01-06_*.py`; this report’s figures are produced by `07_plot_training_curves.py`.
- Trained checkpoint: `data/checkpoint.best.pt` (~14 MB; gitignored, regeneratable in 25 min on one L4).
- Training-history and test-result artifacts: `data/{training_history.json, test_result.json, training_split.parquet}`.
- NeuraLens catalog source: Google Drive CSV (URL above), licensed for non-commercial academic use per the NeuraLens project page.

- DESI DR1 zcatalog source: <https://data.desi.lbl.gov/public/dr1/spectro/redux/iron/zcatalog/v1/zall-pix-iron.fits> (22.4 GB, public).
- DECaLS cutouts: pulled at runtime from legacysurvey.org/viewer/fits-cutout; cutouts gitignored.
- Python environment: `/raid/benson/.venvs/lensfinder — torch 2.12+cu126, torchvision 0.27, astropy, fitsio, scikit-learn, matplotlib`.
- Code repository: <https://github.com/usfcs-edu/agentik-lensing>.

9 Conclusions

We have implemented and trained a PyTorch port of the Lanusse et al. (2018) RESNET-46 under the Huang et al. (2020) hyperparameter recipe, reaching a test AUC of 0.9991 — narrowly above the published 0.98. The architecture, the loss, the optimiser settings, the data augmentation, and the cutout size are all faithful to the published descriptions. Once the methodological substitutions of §7 are corrected, this result is the basis of a faithful Phase 3b reproduction of the full DECaLS deployment that produced Huang et al. (2020)’s 335 new strong-lens candidates.

A Per-epoch training history (compact)

Selected milestones from `data/training_history.json`. First learning-rate decay $1e-3 \rightarrow 1e-4$ at epoch 41; second decay $1e-4 \rightarrow 1e-5$ at epoch 81. Best validation AUC highlighted.

Table 3: Selected milestones from the 120-epoch training run.

epoch	LR	train loss	val loss	val AUC
1	1e-3	0.4158	0.4458	0.5838
4	1e-3	0.2134	0.2159	0.9567
16	1e-3	0.1479	0.2365	0.9808
20	1e-3	0.1179	0.1169	0.9933
34	1e-3	0.0629	0.0499	0.9974
41	1e-4	0.0458	0.1772	0.9926
51	1e-4	0.0348	0.0474	0.9976
71	1e-4	0.0318	0.0405	0.9982
80	1e-4	0.0267	0.0432	0.9983
81	1e-5	0.0177	0.0626	0.9962
88	1e-5	0.0232	0.0479	0.9983
93	1e-5	0.0223	0.0421	0.9983
120	1e-5	0.0178	0.0849	0.9954
test (epoch 93 checkpoint)				0.9991

References

K. He, X. Zhang, S. Ren, and J. Sun. Identity Mappings in Deep Residual Networks. In *European Conference on Computer Vision (ECCV)*, 2016. doi: 10.1007/978-3-319-46493-0_38.

- Y.-M. Hsu, X. Huang, C. Storfer, J. C. Inchausti, et al. A New Way to Discover Strong Gravitational Lenses: Pair-wise Spectroscopic Search from DESI DR1. *arXiv preprint*, 2025.
- X. Huang, C. Storfer, V. Ravi, A. Pilon, et al. Finding Strong Gravitational Lenses in the DESI DECam Legacy Survey. *The Astrophysical Journal*, 894(1):78, 2020. doi: 10.3847/1538-4357/ab7ffb.
- X. Huang, C. Storfer, A. Gu, et al. Discovering New Strong Gravitational Lenses in the DESI Legacy Imaging Surveys. *The Astrophysical Journal*, 909(1):27, 2021. doi: 10.3847/1538-4357/abd62b.
- X. Huang, S. Baltasar, N. Ratier-Werbin, et al. DESI Strong Lens Foundry I: HST Observations and Modeling with GIGA-Lens. *arXiv preprint*, 2025.
- D. P. Kingma and J. Ba. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- F. Lanusse, Q. Ma, N. Li, T. E. Collett, C.-L. Li, S. Ravanbakhsh, R. Mandelbaum, and B. Póczos. CMU DeepLens: Deep Learning For Automatic Image-based Galaxy-Galaxy Strong Lens Finding. *Monthly Notices of the Royal Astronomical Society*, 473(3):3895–3906, 2018. doi: 10.1093/mnras/stx1665.
- L. A. Moustakas et al. The Master Lens Database. <http://masterlens.astro.utah.edu/>, 2012.
- A. Paszke, S. Gross, F. Massa, et al. PyTorch: An Imperative Style, High-Performance Deep Learning Library. In *Advances in Neural Information Processing Systems*, volume 32, 2019.